

DevOps Exam Assignment

May 2024

In this exam assignment you will design and implement an infrastructure and development environment for a web application that enables users to search for contact information of teachers using their 'short name'.

The Vue web application connects to a backend written in Ruby. The contact information is stored in this backend application. The Ruby application uses an SQLite database to store the contacts during development and a MySQL database when the application runs on the live website.

Your task will be to run and test this application. Additionally, you will have to set up the continuous integration pipeline for this application and containerize the application.



Entry Requirements

- It is mandatory to hand in all the *practice assignments* and have them approved by your teacher, before your final assignment can be handed in. If you haven't done so, please contact your teacher ASAP.
- For each requirement (see below) you should explain how you solved the problem and why you solved it in this way. This should be described in your own words in a way that is understandable for the teacher that is grading your work. You can either write it down in a Markdown file included in your repository, or in comments in the appropriate source file (in the latter case, please include references to the locations of these comments in the Markdown file).
- Make sure you use a repository in the Saxion Gitlab organization, created via <https://repo.hboictlab.nl>
- Upload a zipped export of your Gitlab project to Blackboard (Settings → General → Advanced → Export Project).
Please mention the URL of your repository in the comment box on Blackboard as well.

Exam Rules

- The deadline for handing in the assignment is Monday 3rd of June 2024 at 9:00 (week 4.5).
- The assignment is done in exam groups of two students that are registered on Blackboard.
- Your grade is determined based on the oral exam that will take place after you handed in your assignment. Make sure you understand your design well, we'll ask theoretical questions as well.
- You will only receive points for elements of your work that you are able to explain during the oral exam. For example: if your code works but you cannot explain what it does or why it works like that, that part of your solution won't count towards your final grade.
- It is **not** allowed to receive assistance from anyone outside of your exam group, except for your teacher. It is also **not** allowed to give assistance to anyone else.
- You are allowed to use online posts, articles, tutorials, books, videos. You must add references¹ for all sources and code fragments that you used in your text/code.

¹ <https://libguides.murdoch.edu.au/IEEE>

Process Requirements

While doing the assignment you should follow a proper software development process, as was covered in week 1. This includes things like tracking tasks using Gitlab issues, using Git for collaborating on the code, and using branches and merge requests to facilitate code reviews. This should all be evident from your Gitlab repository.

Functional Requirements

1. Add data to the backend

In the template folder you will find a folder containing a backend application. This is a *Ruby On Rails* application provided by us that returns data in JSON format. If you run the application (using NPM) and navigate to <http://localhost:3000/contacts/> you will get a list of contacts (which should be empty at this point).

In order to add contacts to the backend you should create a POST request with the proper format (see `add_contacts.sh`). You can use `curl` to create such a POST request. Make sure you set the content type to `application/json` in your Curl command. The REST endpoint that is used should be configurable using an environment variable.

In the folder '*seed-application*' you will find a Bash script called `add_contacts.sh`. Edit this script so that it sends a list of contacts to the backend. The list should be read from the `data.txt` file. Make sure you implement proper error handling.

2. Containerization of the applications

The backend and frontend application should be containerized using Docker and Docker Compose. There should be separate backend and frontend containers.

You should facilitate a convenient way to build, run and stop the containers by writing a script. Make sure that intermediate files created by running your application outside of Docker are not included in the image.

3. Containerization of an external database and seed script

Extend your compose file so that the data is not stored in SQLite in the backend but in an external MySQL database (which is ran in a container). For this to work you should set the following environment variables in the backend and the MYSQL container:

```
MYSQL_USER
MYSQL_PASSWORD
MYSQL_ROOTPASSWORD
MYSQL_DATABASE
```

In the backend container you should also set the following environment variables:

- `RAILS_ENV` (set its value to 'production' so that Rails knows that it needs to run in production mode)
- `MYSQL_HOST` (set the name of the server which the backend needs to connect to)

To make sure the application is working properly, it is convenient to create an additional container for filling the database with the script that you have written in step 1. Make sure this script is only executed when the database is empty.

4. Continuous Integration

Set up continuous integration for the app. The continuous integration should consist of at least two stages. In the first stage you have to [lint](#) the frontend and backend applications. In the second stage you need to run the provided (and supplemented by you) unit tests. You might have to fix some linting issues first before the linting is successful! You can find the unit tests at `backend/test/controllers/contacts_controller_test.rb`

Make sure that the pipeline is automatically executed when changes are pushed to the main branch or to a merge request. Both the linting and the unit testing will create a report. Make sure to store these reports as artifacts in Gitlab (with an expiration date of one month).

5. Create an infrastructure to run the application

Create an appropriate infrastructure for the application in AWS using Terraform. The infrastructure does not need to be high availability (yet). If you do not manage to do so using Terraform you are allowed to manually create the infrastructure (this will lose you points of course). Your Terraform configuration should be placed in the provided `infra` folder.

The backend should be installed on an EC2 instance (manually for now), the frontend may be served as a static website.

6. Automate the application deployment

Automate the deployment of the backend, so that when we commit to the main branch the Gitlab CI:

- creates the necessary Docker image(s)
- uploads the image(s) to the [private container registry](#) for your Gitlab repository
- connects to the EC2 instance using SSH, pulls the Docker image(s) and runs the image(s).

To successfully achieve this, you should copy the contents of your AWS private SSH key (that you use to connect to the instance) to an environment variable in Gitlab CI (Settings -> CI/CD -> Variables, click on "Expand"). In your `.gitlab-ci.yml` file you can use this variable to connect to the EC2 instance using SSH.

When using SSH the client will check whether the server is a known server. For our automated deployment this is not very useful (as it would require someone to type in 'yes'). To omit this check, you can pass the `-o StrictHostKeyChecking=no` option to SSH.

7. Create a high availability backend

Extend your Terraform configuration with a high availability backend. Don't forget to make the database highly available (i.e. redundant) as well.

Make sure the instances are not accessible directly from outside of your network. Remember to update the frontend to connect to your updated infrastructure.